

# Harneet Kaur

## Data Scientist | LLM & Agentic AI Applications

+91 8018724774 | harneetkaur4464@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

### PROFILE

Data Scientist (M.Sc. Data Science, 9.46 CGPA) building Agentic AI, RAG, and LLM applications. Shipped a 6-agent LangGraph framework for clinical AI routing and production automation-engine code at Incredible Visibility. Statistics-to-data-science background with full-stack engineering fluency.

### WORK EXPERIENCE

**Full-Stack Engineering Intern — AI Automation Platform** / *Incredible Visibility* / *Remote* / *Mar 2026 – Present*

**Automation Workflow Engine (Jul 2026 – Present) — n8n-style visual workflow platform**

- Built backend runner and validation for a Switch (n-way branch) node using the new `executeContract(ctx)` pattern — now the reference implementation for new node types.
- Shipped builder UI (canvas, config panels, registration) for 5 node types — Switch, Loop, Merge, Parallel, Sub-workflow — enabling visual branching and iterative flows.
- Built an expression/variable picker and run-state inspector across all builder fields; fixed variable resolution bugs affecting 9 of 13 node types.
- Delivered 2 end-to-end golden-flow templates (scheduled reports, webhook reconciliation); resolved cross-org data leakage and a prompt-injection risk flagged in security review.

**Platform & Product Engineering (Mar 2026 – Jul 2026)**

- Built a full-stack, entity-agnostic fine-tuning pipeline for Agents/Assistants — CSV-to-JSONL normalization, dual Supabase/OpenAI storage, polling-based status tracking.
- Led V2 UI migrations for Settings, Profile, and Tools Library — 60+ files across frontend, API, and Supabase.
- Resolved 25+ UI/UX and theming defects across dashboard, builder, and pricing surfaces.
- Shipped 11+ merged PRs across backend, frontend, and database layers.

### ACADEMIC PROJECTS

**SEHAJ: Multi-Agent Conversational AI for Mental Health Support** / *LangGraph, Llama 3.3, Groq, Streamlit*

- Engineered a multi-agent orchestration framework with 6 specialized agents using LangGraph, improving intent routing and reducing context interference in clinical AI workflows.
- Designed a Guardian safety protocol achieving 90% interception accuracy on high-risk scenarios with 0% false positives on benign emotional conversations.
- Reduced token overhead by 54.7% and achieved 0.71s average latency via Groq acceleration; attained a CTRS therapeutic validation score of 4.47/5.

**NLP-Powered Mental Health Status Detection** / *Python, Transformers, Scikit-Learn*

- Built an NLP pipeline classifying 51,000+ social media posts into 7 mental health categories, addressing severe class imbalance through evaluation-driven modeling.
- Benchmarked MentalBERT, BioBERT, and PubMedBERT; MentalBERT performed best (F1 = 0.84).
- Improved baseline performance by 15% through emotion-aware preprocessing; deployed via Streamlit.

### TECHNICAL SKILLS

**Machine Learning & AI:** Statistics, Probability, Machine Learning, NLP, Transformers, LLM Applications, Agentic AI, LangGraph, Fine-Tuning, Model Evaluation, Scikit-Learn, TensorFlow, Keras, RAG, Prompt Engineering, LLM Provider Architecture

**Programming:** Python, SQL, R, JavaScript

**Full-Stack & Platform:** React, Node.js, Supabase, REST APIs, Git, Bitbucket

**Data & Visualization:** Pandas, NumPy, Matplotlib, Seaborn, Streamlit, Jupyter Notebook

### EDUCATION

- VIT University, Vellore | M.Sc. Data Science | 2024–2026 | CGPA: 9.46/10
- BJB Autonomous College, Bhubaneswar | B.Sc. (Hons.) Statistics | 2020–2023 | CGPA: 9.12/10, Rank 2 in Department (Distinction)
- AIR 607 (IIT JAM Statistics) | AIR 348 (CUET PG)