Harneet Kaur

Data Scientist | LLM & Agentic AI Applications

+91 8018724774 | harneetkaur4464@gmail.com | [LinkedIn] | [GitHub] | [Portfolio]

PROFILE

Data Scientist specializing in Agentic AI, RAG, and LLM applications, built on an M.Sc. in Data Science. Built a 6-agent LangGraph framework for clinical AI routing during the degree, carried into industry experience as a Full-Stack Engineering Intern at Incredible Visibility. Brings a rare combination: the statistical rigor to build models that hold up, and the engineering discipline to ship them into real systems. Looking for full-time AI/ML Engineer roles.

WORK EXPERIENCE

Full-Stack Engineering Intern, AI Automation Platform | *Incredible Visibility* | *Remote* | *Mar 2026 – Present*

Platform Architecture Design

- Authored the architectural outline for AI Agent Triggers, Automations & Intelligent Decision Engine, separating the platform into a deterministic workflow orchestrator and an isolated, schema-validated AI Agent.
- Designed the trigger normalization layer (webhook, manual, schedule) into a single event bus, and defined the bounded ReAct execution loop and guardrails for the AI Agent node.
- Specified the credential vault, tool response envelopes, and node contract interface; the document became the reference architecture for the 23-ticket Automation Workflow Engine v2 epic.

Automation Workflow Engine (July 2026 – Present), n8n-style visual workflow platform

- Contributed to the builder UI and logic for 5 node types (Switch, Loop, Merge, Parallel, Sub-workflow), enabling visual branching and iterative flows.
- Worked on an expression/variable picker and run-state inspector, debugging variable resolution issues across 9 of 13 node types.
- Contributed to 2 end-to-end golden-flow templates (scheduled reports, webhook reconciliation); helped catch and resolve a cross-org data leakage issue and a prompt-injection risk flagged in security review.

Platform & Product Engineering (Mar 2026 – July 2026)

- Contributed to a full-stack, entity-agnostic fine-tuning pipeline for Agents/Assistants: CSV-to-JSONL normalization, dual Supabase/OpenAI storage, polling-based status tracking.
- Worked across V2 UI migrations for Settings, Profile, and Tools Library, spanning 60+ files across frontend, API, and Supabase.
- Debugged and refined 25+ UI/UX and theming defects across dashboard, builder, and pricing surfaces.
- Shipped changes for 11+ merged PRs across backend, frontend, and database layers.

ACADEMIC PROJECTS

SEHAJ: Multi-Agent Conversational AI for Mental Health Support | *LangGraph, Llama 3.3, Groq, Streamlit* **(Dec 2025 – May 2026)**

- Engineered a multi-agent orchestration framework with 6 specialized agents using LangGraph, improving intent routing and reducing context interference in clinical AI workflows.
- Designed a Guardian safety protocol achieving 90% interception accuracy on high-risk scenarios with 0% false positives on benign emotional conversations.
- Reduced token overhead by 54.7% and achieved 0.71s average latency via Groq acceleration; attained a CTRS therapeutic validation score of 4.47/5.

NLP-Powered Mental Health Status Detection | *Python, Transformers, Scikit-Learn* **(July 2025 – Nov 2025)**

- Built an NLP pipeline classifying 51,000+ social media posts into 7 mental health categories, addressing severe class imbalance through evaluation-driven modeling.
- Benchmarked MentalBERT, BioBERT, and PubMedBERT; MentalBERT performed best (F1 = 0.84).
- Improved baseline performance by 15% through emotion-aware preprocessing; deployed via Streamlit.

Hybrid Product Recommendation System | *Sentence-BERT, SVD, Surprise, Scikit-learn (co-authored, VIT M. Sc.)* ***(Aug 2024 - Oct 2025)***

- Built a hybrid recommender combining SVD-based collaborative filtering (GridSearchCV-tuned, RMSE 0.97–0.99) with Sentence-BERT semantic embeddings (384-dim) for content-based filtering, addressing cold-start and data sparsity in a 73K-review Amazon dataset.
- Achieved Precision@10 of 13.3–13.8%, Recall@10 of 32–40%, NDCG@10 of 24–29%, outperforming a review-sentiment-only baseline (Elahi et al., 2023) by ~4x on Precision@10.

TECHNICAL SKILLS

Machine Learning & AI: RAG, NLP, LLM Applications, LLM Provider Architecture, Prompt Engineering, Agentic AI, LangGraph, Transformers, Fine-Tuning, Machine Learning, Model Evaluation, Scikit-Learn, TensorFlow, Keras, Probability, Statistics

Programming: Python, SQL, R, JavaScript

Full-Stack & Platform: React, Node.js, Supabase, REST APIs, Git, Bitbucket

Data & Visualization: Pandas, NumPy, Matplotlib, Seaborn, Streamlit, Jupyter Notebook

EDUCATION

- VIT University, Vellore | M.Sc. Data Science | 2024–2026 | CGPA: 9.46/10
- BJB Autonomous College, Bhubaneswar | B.Sc. (Hons.) Statistics | 2020–2023 | CGPA: 9.12/10 (Core CGPA: 9.29), Rank 2 in Department (Distinction)
- AIR 607 (IIT JAM Statistics) | AIR 348 (CUET PG)